

Coding Bootcamp Code in Python

PYTHON FOR SCIENTIFIC COMPUTING: NUMPY

Out of the box

- Python is interpreted
 - Python is slow
 - Python is really slow
- Okay for one-offs, prototypes, short runtimes
- ***Not okay*** for computationally intensive tasks!

Don't use vanilla Python for computations!!!

Python performance

500 × 500 matrices

Python: 0.09 s

C: 0.014 s

Fortran: 0.012 s

```
def init_matrix(n):
```

```
    m = []
```

```
    for i in range(n):
```

```
        m.append([])
```

```
        for j in range(n):
```

```
            m[i].append(random.random())
```

```
    return m
```

Represent matrix as list of lists

```
def matmul(a, b, c):
```

```
    n = len(a)
```

```
    for i in range(n):
```

```
        for j in range(n):
```

```
            c[i][j] = 0.0
```

```
            for k in range(n):
```

```
                c[i][j] += a[i][k]*b[k][j]
```

$$C = A \cdot B$$

$$C_{ij} = A_{i1}B_{1j} + \dots + A_{iN}B_{Nj}$$

Python: 32 s

C: 0.49 s

Fortran: 0.11 s

Naive!

Libraries for numeric computation

- `numpy`
 - Fast arrays
 - Matrix operations (BLAS-like)
 - Linear algebra
 - Fast Fourier Transform
 - Mathematical functions defined on arrays
 - Pseudo-random number generation to initialize arrays
 - Simple statistics
- ...

Python using numpy

500 × 500 matrices

numpy: 0.011 s

numpy: 0.077 s


```
import numpy as np

def init_matrix(n):
    return np.random.uniform(0.0, 1.0, (n, n))

def matmul(a, b):
    return a @ b
```

Language/library	Execution time matrix multiplication (s)
Python	32
C	0.49
Fortran	0.11
Python/numpy	0.077
Fortran/BLAS	0.060

415 ×



Creating array I

convention

```
import numpy as np
```

create 3-element array, all 0.0

```
v1 = np.zeros(3)
```

```
v2 = np.ones(3)
```

create 3-element array, all 1.0

```
v3 = np.empty(3)
```

create 3-element "empty" array
initialize elements later

create 2×3 array, all elements 0.0

```
a = np.zeros((2, 3))
```

```
b = np.ones((2, 3))
```

create 2×3 array, all elements 1.0

```
c = np.eye(2)
```

create 2×2 identity array


```
d = np.empty((2, 3))
```

create 2×3 "empty" array
initialize elements later

Default type: `np.float` \equiv double precision


Creating arrays II

```
e = np.random.uniform(0.0, 1.0, (2, 3))
```




create 2×3 array, elements x randomly drawn from uniform distribution such that $x \in [0.0, 1.0[$

```
f = np.array([[3.1, 4.2, -1.1], [-0.3, 1.3, 13.1]])
```



create 2×3 array from a Python list of lists

```
f = np.genfromtxt('matrix.txt')
```




create 2×3 array from text file

1.2	2.3	3.4
4.5	5.6	6.7

Creating arrays III


```
e = np.arange(-1.0, 1.0, 0.25)
```



create 8-element array, first element -1.0,
last element less than 1.0, step 0.25

```
[ -1.  -0.75 -0.5  -0.25  0.   0.25  0.5   0.75 ]
```

```
f = np.linspace(-1.0, 1.0, 9)
```



create 9-element array, first element -1.0,
last element 1.0, determine step

```
[ -1.  -0.75 -0.5  -0.25  0.   0.25  0.5   0.75  1. ]
```


Numpy data types

- Integers

`np.int8, np.int16, np.int32, np.int64`
default for `np.int`: `np.int32` on 32-bit, `np.int64` on 64-bit architecture (`np.uint<n>` for unsigned integers)

- Floating point numbers

`np.float16, np.float32, np.float64, np.float96`
default for `np.float`: `np.float64`, i.e., double precision

- Complex numbers

`np.complex64, np.complex128, np.complex192`
default for `np.complex`: `np.complex128`, i.e., double precision

- Boolean values: `np.bool`

- Characters/`np.char`

```
v = np.zeros(3, dtype=np.int8)
```

Accessing array elements

```
a = np.zeros((2, 3))
```

- Array dimensions, strides

```
a.shape == (2, 3)
```

```
a.strides == (24, 8)
```

nr. bytes to
next row

nr. bytes to
next column

- Assigning to a specific element

```
a[1, 0] = 5.0
```

0.0	0.0	0.0
5.0	0.0	0.0

- Using an element's value

```
q = a[1, 0] + a[1, 2]
```

Note:

- Implicit conversion of tuple for indexing
- 0-based indexing

Accessing subarrays: slicing

```
a = np.arange(1, 21).reshape(4, 5)
```

- Second column

```
a[:, 1]
```

[1	2	3	4	5]
[6	7	8	9	10]
[11	12	13	14	15]
[16	17	18	19	20]

[2	7	12	17]
---	---	---	----	----	---

- Third row

result: 1D array

```
a[2, :]
```

[11	12	13	14	15]
---	----	----	----	----	----	---

- 2D subarray

```
a[1:3, 1:3] = np.eye(2)
```

[1	2	3	4	5]
[6	1	0	9	10]
[11	0	1	14	15]
[16	17	18	19	20]

cfr. list slicing, but...
array slicing does **not** copy!

Fancy indexing

```
a = np.arange(1, 10, dtype=np.int).reshape(3, 3)
```

```
[[ 1  2  3 ]  
 [ 4  5  6 ]  
 [ 7  8  9 ]
```

- Conditional indexing

```
a[a % 2 == 1] = 0
```

3 × 3 Boolean array



```
[[ 0  2  0 ]  
 [ 4  0  6 ]  
 [ 0  8  0 ]
```

- Conditional assignment

```
np.where(a > 0, 1, -1)
```

3 × 3 Boolean array



```
[[ -1  1 -1 ]  
 [ 1 -1  1 ]  
 [ -1  1 -1 ]
```

Operations on arrays

- Scalar-array operations: $+$, $-$, $*$, $/$, $//$, $**$

```
a = np.array([[1.0, 3.0], [4.0, 5.0]])  
print(3.0 + a)
```

```
[[ 4.  6. ]  
 [ 7.  8. ]]
```

- Element-wise operations: $+$, $-$, $*$, $/$, $//$, $**$

```
a = np.array([[1.0, 3.0], [4.0, 5.0]])  
b = np.array([[2.0, 3.0], [1.0, 0.5]])  
print(a*b)
```

```
[[ 2.  9. ]  
 [ 4.  2.5 ]]
```

- Matrix product

```
print(np.dot(a, b))
```

```
[[ 5.  4.5 ]  
 [ 13. 14.5 ]]
```

– Python 3.5 style

```
print(a @ b)
```

Functions operating on arrays

```
a = np.empty((500, 500))  
b = np.random.uniform(0.0, 1.0, (500, 500))
```

- **Avoid**

```
for i in range(500):  
    for j in range(500):  
        a[i, j] = math.sqrt(b[i, j])
```

140 μ s

- **Use**

```
a = np.sqrt(b)
```

6.5 μ s

21 \times

- **Other functions:** `np.sin, ..., np.sinh, ..., np.exp, np.trace, np.transpose, ...`

Some linear algebra

```
a = np.array([[1.0, 3.0], [4.0, 5.0]])
```

- numpy has some linear algebra operations

- matrix power

```
np.linalg.matrix_power(a, 3)
```

```
[[ 85. 129.]  
 [172. 257.]
```

- matrix inverse

```
np.linalg.inv(a)
```

```
[[ -0.71428571  0.42857143]  
 [ 0.57142857 -0.14285714]]
```

- determinant

```
np.linalg.det(a)
```

```
-7.0
```

- eigen values

```
np.linalg.eigvals(a)
```

```
[-1.  7.]
```

References versus copies

- Reshape: different view on same data

```
a = np.array([[1.0, 3.0],  
              [4.0, 5.0]])
```

```
[[ 1.  3.]  
 [ 4.  5.]
```

```
b = a.reshape((a.size, ))
```

```
[ 1.  3.  4.  5.]
```

```
a[0, 0] = 13.0
```

a

```
[[ 13.  3.]  
 [  4.  5.]
```

b

```
[ 13.  3.  4.  5.]
```

- Some operations return copies, check documentation carefully

numpy data I/O revisited

- Reading text file with 10^9 64-bit floats

- `np.loadtxt(...)`: 57 minutes,

44 GB RAM

- `np.fromfile(..., sep='\\n')`: 4.6 minutes,

8 GB RAM

5 ×

12 ×

All functions are equal, but...
some are more equal than others

35 ×

- Reading binary file with 10^9 64-bit floats

- `np.fromfile(...)`: 8 seconds,
8 GB RAM

Not all data formats are equal: HDF5 to the rescue

Matrices

- Matlab-like initialization

```
a = np.matrix('1.0  3.0; 4.0  5.0')
```

```
[[ 1.  3. ]  
 [ 4.  5. ]]
```

- Overloaded `*` and `**` operators

```
a = np.matrix([[1.0, 3.0], [4.0, 5.0]])  
b = np.matrix([[2.0, 3.0], [1.0, 0.5]])  
print(a*b)  
print(a**3)
```

```
[[ 5.  4.5 ]  
 [ 13. 14.5 ]]
```

```
[[ 85. 129. ]  
 [ 172. 257. ]]
```

- Result is always matrix (2D)

```
a = np.matrix('1.0, 3.0')  
b = np.matrix('2.0; 4.0')  
print(a*b)
```

```
[[ 14. ]]
```

Don't use,
confusing

References

- numpy for MATLAB users

<http://mathesaurus.sourceforge.net/matlab-numpy.html>

Code Pack 25

A. Numpy: Make the code

Coding Bootcamp Code in Python

PYTHON FOR SCIENTIFIC COMPUTING: SCIPY

Libraries for numeric computation

- ...
- `scipy`
 - Dense/sparse linear algebra
 - Solving ordinary differential equations
 - Numerical integration
 - Optimization
 - Interpolation
 - Signal processing
 - Statistics
 - Special mathematical functions
 - Mathematical & physical constants

- ...

```
import scipy as sp
```

convention



```
import scipy.linalg
```

import subpackages as needed



Singular Value Decomposition

Should be fast when built against good BLAS/Lapack library

- Computing SVD

```
import scipy.linalg
a = np.array([[7.3, 5.7], [-1.2, 5.3]])
u, s, v = sp.linalg.svd(a)
```

- Note: s is not a 2D-array, it is a 1D-array

```
S = np.diag(s)
```

- Let's check

```
A = u @ S @ v
delta = A - a
```

```
[[ 8.88178420e-16,  0.00000000e+00],
 [ 4.44089210e-16,  0.00000000e+00]]
```

Linear regression

- Reading data

```
x, y
0.000e+00, 1.206e+00
5.263e-02, 1.207e+00
...
data.csv
```

```
data = np.genfromtxt('data.csv',
                     dtype=[np.float64, np.float64],
                     delimiter=',', names=True)
```

- Linear regression

```
import scipy as sp
import scipy.stats
slp, intc, r, _, _ = sp.stats.linregress(data['x'],
                                         data['y'])
```


Optimization: function definitions

- Minimize $f(x, y) = (x^2 + y^2)^2 - 2x^2 - 2y^2 + 0.1x$
- Define function

```
def f(X):  
    x = X[0]  
    y = X[1]  
    return (x**2 + y**2)**2 - 2*x**2 - 2*y**2 + 0.1*x
```

- Define gradient

```
def grad_f(X):  
    x = X[0]  
    y = X[1]  
    f_x = 4*(x**2 + y**2)*x - 4*x + 0.1  
    f_y = 4*(x**2 + y**2)*y - 4*y  
    return np.array([f_x, f_y])
```

Optimization

- Compute minimum

```
import scipy.optimize

x0 = np.array([1.0, 0.01])
xopt = scipy.optimize.fmin_cg(f, x0, fprime=grad_f,
                             disp=False)
```

- Many methods
 - Powell
 - Conjugate gradient
 - BFGS
 - Newton conjugate gradient
 - ...

Ordinary differential equations

- Rewrite higher order differential equation to set of first order equations

$$\frac{d^2\theta}{dt^2} = -\frac{g}{l}\theta - q\omega + F_D \sin \Omega_D t \quad \Leftrightarrow \quad \begin{cases} \frac{d\theta}{dt} = \omega \\ \frac{d\omega}{dt} = -\frac{g}{l}\theta - q\omega + F_D \sin \Omega_D t \end{cases}$$

```
def func(t, y, g, l, q, F_D, Omega_D):  
    return [  
        y[1],  
        -(g/l)*y[0] - q*y[1] + F_D*np.sin(Omega_D*t)  
    ]
```

$$y[0] \equiv \theta \quad y[1] \equiv \omega$$

Jacobian for equations

- For many methods, convergence improves by specifying Jacobian


$$\begin{cases} f_1(\theta, \omega, t) = \omega \\ f_2(\theta, \omega, t) = -\frac{g}{l}\theta - q\omega + F_D \sin \Omega_D t \end{cases} \quad \begin{bmatrix} \frac{\partial f_1}{\partial \theta} & \frac{\partial f_1}{\partial \omega} \\ \frac{\partial f_2}{\partial \theta} & \frac{\partial f_2}{\partial \omega} \end{bmatrix}$$

```
def jac(t, y, g, l, q, F_D, Omega_D):  
    return [  
        [0.0, 1.0],  
        [-g/l, -q]  
    ]
```

Integrate ODEs

- Integrate from t_0 to t_{\max} in steps Δt

integration method
RKF(4, 5)



```
from scipy.integrate import ode
...
ode_sys = ode(func, jac).set_integrator('dopri5')
ode_sys.set_initial_value([theta0, omega0], t0)
ode_sys.set_f_params(g, l, q, F_D, Omega_D)
ode_sys.set_jac_params(g, l, q, F_D, Omega_D)

while ode_sys.successful() and ode_sys.t < t_max:
    ode_sys.integrate(ode_sys.t + delta_t)
    print(ode_sys.t, ode_sys.y[0], ode_sys.y[1])
```

Signal processing

- Remove noise from sound file (WAV)
 - Read WAV file



```
...  
from scipy.io import wavfile  
...  
sample_rate, signal = wavfile.read(wav_file_name)
```

- Perform FFT to compute frequency spectrum

```
...  
n = len(signal)  
freq = sample_rate*np.arange(n)/n  
Y = sp.fft(signal)/n
```



Original signal



Create highpass filter

- Import signal processing package

```
...  
import scipy.signal  
...
```

- Create IIR digital filter

```
...  
b, a = sp.signal.iirfilter(17, cutoff,  
                           rs=min_attenuation,  
                           btype='highpass',  
                           analog=False,  
                           ftype='cheby2')  
...
```

The diagram shows five callout boxes with arrows pointing to specific arguments in the `iirfilter` function call:

- order of the filter**: Points to the value `17`.
- fraction of Nyquist frequency**: Points to the variable `cutoff`.
- minimum attenuation**: Points to the variable `rs=min_attenuation`.
- filter type**: Points to the string `btype='highpass'`.
- IIR filter type**: Points to the string `ftype='cheby2'`.

Filter signal

- Apply filter



```
...  
base = np.uint8(np.mean(signal))  
filtered_signal = sp.signal.filtfilt(b, a, signal)  
wav_signal = base + np.array(filtered_signal,  
                             dtype=np.uint8)  
...
```

- Write signal to WAV file

```
...  
wavfile.write(filtered_wav_file_name, sample_rate,  
              wav_signal)  
...
```



Filtered signal



Code Pack 26

A. Scipy: See the code